

Åbo Akademi University

Systematic Mapping and Benchmarking of Sign Language Datasets

Author: Kaniz Fatema Begum

Supervisor: Luigia Petre

Content

[1. Introduction](#)

[2. Background and Related Work](#)

[2.1 Sign Language Datasets in Computational Research](#)

[2.2 Types of Sign Language Datasets](#)

[2.3 Data Modalities](#)

[2.4 Annotation in Sign Language Datasets](#)

[2.5 Dataset Scale and Signer Diversity](#)

[2.6 Benchmarking and Evaluation Protocols](#)

[2.7 Dataset Accessibility and Licensing](#)

[2.8 Related Work](#)

[2.8.1 Surveys of Sign Language Recognition and Translation](#)

[2.8.2 Studies Directly Using Sign Language Datasets](#)

[2.8.3 Systematic Mapping Methodology](#)

[2.8.4 Positioning This Thesis](#)

[4. Results and Initial Discussion](#)

[4.1 Overview of Included Datasets](#)

[4.2 Distribution by Sign Language and Region](#)

[4.3 Dataset Type and Task Support](#)

[4.4 Dataset Scale and Composition](#)

[4.5 Modalities and Recording Setup](#)

[4.6 Annotation and Benchmarking Practices](#)

[4.7 Accessibility and Availability](#)

[4.8 Initial Discussion](#)

[5. Conclusion](#)

[References](#)

1. Introduction

Sign languages are natural languages used by deaf and hard-of-hearing communities worldwide. In contrast to spoken languages, sign languages are visually conveyed through hand movements, facial expressions, and body posture, and have their own distinct linguistic structures. Recent advances in artificial intelligence and machine learning have increased research interest in automatic sign language recognition, translation, and accessibility technologies. Progress in these domains relies on the availability of suitable and well-documented datasets.

In the past decade, a wide range of sign language datasets have been introduced, covering various sign languages, recording environments, and research objectives. These datasets vary significantly in scale, modality (such as RGB video or depth data), annotation schemes, and signer diversity. However, information regarding these datasets is scattered across academic publications, project websites, and repositories. Consequently, researchers, particularly those new to the field, frequently encounter challenges in identifying suitable datasets and comprehending their limitations.

The lack of a consolidated and structured overview of sign language datasets hinders research progress and reproducibility. In the absence of clear comparisons and classifications, it is difficult to assess dataset usability, identify gaps in language coverage, or determine the suitability of datasets for specific research tasks. This challenge highlights the necessity for a systematic approach to mapping and evaluating existing sign language datasets.

This thesis aims to deliver a structured overview of existing sign language datasets by identifying, categorizing, and comparing their key characteristics. The analysis focuses on properties such as language coverage, modality, scale, annotation, and accessibility. By applying a systematic mapping methodology, the thesis clarifies the current landscape of sign language datasets and highlights existing limitations and research gaps.

2. Background and Related Work

This chapter establishes the conceptual and scholarly context for the systematic mapping study presented in this thesis. It outlines the key concepts, prior research, and methodological frameworks that inform the study's design and scope.

2.1 Sign Language Datasets in Computational Research

Sign language datasets serve as the primary empirical foundation for computational sign language research (Koller, 2020; Rastgoo et al., 2021). These datasets comprise video recordings of one or more signers, with each clip annotated to indicate the identity of the sign or signs being performed. They are utilized to train, evaluate, and benchmark models for a variety of tasks (Koller, 2020).

Existing datasets differ substantially in scope, including the sign language represented, the number of signs and signers, the type of data recorded, and the level of annotation detail (Bragg et al., 2021; Koller, 2020). These variations reflect distinct research priorities and directly influence the tasks a dataset can support and the reliability of experimental results (Koller, 2020). Understanding these characteristics is essential for selecting appropriate datasets and for comparing results across studies (Wadhawan and Kumar, 2021).

2.2 Types of Sign Language Datasets

A primary distinction among sign language datasets is the type of signing captured: isolated signs, continuous signing, or a combination of both.

An isolated sign dataset consists of short video clips in which a single sign is performed independently. Each clip is labelled with the name of the sign, known as a gloss, which is a written word used to identify a sign (e.g., WATER, HELLO). These datasets support sign language recognition (SLR), the task of identifying the sign being performed in a given video clip. A widely used example is WLASL (Word-Level American Sign Language), which contains over 21,000 clips covering 2,000 signs performed by 119 signers.

A continuous signing dataset captures naturally flowing sign language, where multiple signs are produced sequentially to form phrases and sentences. These datasets are more complex to collect and annotate because sign boundaries must be identified within the video stream. They support tasks such as sign-spotting, which involves detecting where specific signs occur within a continuous signing stream, and sign language translation (SLT), which converts a signed video into a written or spoken sentence in another language. An example is

RWTH-PHOENIX-Weather 2014, a German Sign Language dataset of weather forecast broadcasts containing 6,841 signed sentences from 9 signers (Forster et al., 2014).

A third category, referred to as mixed-type datasets, combines isolated signs, continuous sentences, and, in some cases, fingerspelling sequences within a single release. Examples include the SL-ReDu GSL corpus, which provides isolated, continuous, and fingerspelling subsets for Greek Sign Language, and AzSLD, which covers fingerspelling letters, isolated words, and continuous sentences for Azerbaijani Sign Language. Mixed-type datasets are particularly valuable as they enable models to be evaluated across multiple task types using a consistent signer pool.

2.3 Data Modalities

The term modality refers to the type of sensor or data format used to record signing. Datasets may include one or multiple modalities, each capturing distinct aspects of the signing event.

RGB video, which is standard color video from a conventional camera, is the most common modality. It captures the full visual appearance of signing but is sensitive to lighting conditions and background clutter. Depth data, recorded with sensors such as the Microsoft Kinect, provides distance information for each pixel, enabling more robust foreground segmentation and three-dimensional analysis, although it requires specialized hardware.

Skeleton or pose data represents the signer as a set of body joint coordinates, such as wrist, elbow, and finger positions. These coordinates may be estimated automatically from RGB video using tools such as MediaPipe or captured directly via motion-capture equipment. Pose data is compact and less sensitive to visual factors such as clothing; however, automatic estimation can be unreliable for rapid hand movements or when one hand obscures the other.

Some datasets record signing simultaneously from multiple camera angles (multi-view), which reduces occlusion but substantially increases the complexity of data collection.

2.4 Annotation in Sign Language Datasets

An annotation is a descriptive label assigned to raw video data. The depth of annotation determines the range of tasks a dataset can support.

The most common form of annotation is the gloss label, which assigns a sign name to each clip or segment. In continuous datasets, annotations must also specify the start and end points of each sign, a process known as temporal segmentation, which is necessary for sign-spotting

tasks. Datasets intended for SLT additionally provide a spoken-language translation of each signed phrase, pairing signed video with a corresponding written sentence (for example, PHOENIX-2014-T includes German sentence translations alongside gloss labels) (Camgöz et al., 2018). A small number of datasets include phonological annotations that describe the sub-lexical structure of signs, such as hand shape, location, and movement, primarily for linguistic analysis rather than recognition.

2.5 Dataset Scale and Signer Diversity

The scale of a dataset is typically measured by vocabulary size (the number of distinct sign classes), the total number of video samples, and the number of signers. These dimensions collectively determine the extent of variation within the dataset and the degree to which a trained model can generalize to previously unseen signing.

A larger vocabulary and a greater number of samples per sign generally support more robust model training. The number of signers is particularly important; a dataset with only five signers exposes the model to a limited range of signing styles, which can artificially inflate performance when tested on the same individuals. Datasets with more signers, such as AUTSL (43 signers) or WLASL (119 signers), offer greater diversity, although they remain limited compared to the full breadth of a signing community.

Demographic factors such as age, gender, handedness, and whether signers are deaf native users or hearing learners also influence the representativeness of a dataset. However, this information is often absent or inconsistently reported, which complicates the assessment of potential bias in trained models.

Geographic and linguistic concentration represents a further source of imbalance in the existing dataset landscape. Bragg et al. (2021) observe that existing datasets are heavily concentrated in a small number of sign languages, with many of the world’s sign languages remaining without dedicated computational resources. Koller (2020) similarly notes that large-vocabulary continuous recognition benchmarks have historically been limited to a small set of languages, most notably German Sign Language, as seen in RWTH-PHOENIX-Weather.

2.6 Benchmarking and Evaluation Protocols

Benchmarking involves evaluating a model on a standardized test to facilitate comparison across studies. The evaluation protocol, which defines how the dataset is divided into training and test portions, has a significant impact on reported performance and its practical interpretation.

In a signer-dependent evaluation, some signers appear in both the training and test sets, meaning the model has already been exposed to each test signer's style, which can inflate accuracy. In contrast, a signer-independent evaluation ensures that test signers are absent from the training set, providing a more realistic measure of generalization.

Signer-independent protocols are widely regarded as the more rigorous standard but require datasets with a sufficient number of signers to allow for separation across splits. When a fixed train, validation, and test split is published alongside the dataset as a standardized benchmark, all research groups can evaluate under identical conditions, making results directly comparable. In the absence of such a standard, different groups may partition the same data differently, resulting in outcomes that are difficult to compare. Performance is typically reported using metrics such as Top-1 accuracy for isolated recognition or Word Error Rate (WER) for continuous tasks.

2.7 Dataset Accessibility and Licensing

Accessibility refers to whether a dataset is publicly available, whether it requires a formal application, and the licensing terms that govern its use. Many sign language datasets are distributed under research-only licenses; some are open, while others are restricted to academic institutions. A small number have become unavailable over time due to institutional changes or data protection regulations.

The quality of documentation also affects practical usability. A well-documented dataset includes descriptions of the collection procedure, signer demographics, annotation methodology, file formats, and baseline evaluation results. Inconsistent or incomplete documentation hinders the reproducibility of experiments and limits a dataset's adoption within the research community.

2.8 Related Work

2.8.1 Surveys of Sign Language Recognition and Translation

Several surveys have examined computational sign language research, although their primary focus remains on methods and algorithms rather than datasets. Koller (2020) presents a large-scale meta-analysis of sign language recognition, covering approximately 300 published papers and over 400 experimental results from 1983 to 2020. A key finding relevant to this thesis is that a single dataset, RWTH-PHOENIX-Weather, served as the only large-vocabulary continuous sign language recognition benchmark worldwide at the time, illustrating the uneven progression of dataset development across task types.

Sarhan and Frintrop (2023) review isolated sign language recognition methods over the preceding decade, addressing modality types, feature extraction approaches, and transfer learning strategies. While Sarhan and Frintrop (2023) provide a detailed overview of the most prominent ISLR benchmarks, their scope is intentionally restricted to eight high-profile datasets from the last decade and explicitly refrains from including smaller or non-standard datasets. A subsequent review by Tan et al. (2024) expands the scope to include sign language translation and generation, and includes a dedicated analysis of nineteen major datasets, which identifies substantial variability in annotation conventions as a primary obstacle to developing robust, generalizable systems.

Rastgoo et al. (2021) conduct a deep survey of vision-based deep learning approaches to sign language recognition, covering features, modalities, evaluation metrics, and datasets across five years of published research. Datasets are discussed by modality type and scale and used to benchmark reported model performance. While datasets are categorized by basic annotation targets (e.g., hand or face), deeper characteristics such as annotation granularity, signer-independence verification protocols, and the quality of provided documentation are not analyzed as primary comparative dimensions. Wadhawan and Kumar (2021) undertake a decade-long systematic literature review of sign language recognition systems, classifying papers by method type, modality, and application domain. Their study is one of the few in the field to apply a structured review methodology, and it acknowledges variability in dataset properties across studies, but does not treat dataset characteristics as the primary object of structured analysis.

While these surveys collectively establish the importance of dataset metadata such as vocabulary size and signer diversity, their analytical frameworks remain primarily system-centric, which utilizes datasets as benchmarks for model performance. Consequently, there remains a lack of research that treats sign language datasets as the primary object of structured architectural and qualitative analysis.

2.8.2 Studies Directly Using Sign Language Datasets

Several influential papers focus primarily on introducing or utilizing specific sign language datasets, thereby demonstrating the research community's reliance on well-constructed datasets. Li et al. (2020) introduce WLASL, a large-scale American Sign Language dataset containing 2,000 sign classes and recordings from 119 signers, and show that scale and signer diversity significantly influence recognition performance. Sincan and Keles (2020) present AUTSL, a multimodal Turkish Sign Language dataset with both RGB and depth modalities recorded from 43 signers, and provide baseline benchmark results, serving as a concrete example of the benchmarking practices discussed in Section 2.6.

Camgöz et al. (2020) utilize the PHOENIX-2014T dataset to develop a transformer-based model capable of performing both continuous sign recognition and sign language translation within a single framework, demonstrating that annotation depth, specifically the availability of both gloss labels and sentence translations, directly enables more advanced research tasks. Duarte et al. (2021) introduce How2Sign, a large-scale continuous ASL dataset featuring more than 35,000 signed sentences and a 16,000-word vocabulary across multiple modalities, illustrating the scale and diversity necessary to support real-world translation research.

Bragg et al. (2021) adopt a different perspective, examining sign language AI datasets through the lens of Fairness, Accountability, Transparency, and Ethics (FATE). Their analysis reveals that all existing datasets include only a small number of signers, demographic information is rarely reported, and the rights and responsibilities of data contributors are insufficiently addressed. This work most closely aligns with the present thesis in that it subjects sign language datasets, rather than the models trained on them, to critical scrutiny.

2.8.3 Systematic Mapping Methodology

The methodology underpinning this thesis draws on established guidelines for systematic reviews and systematic mapping studies in software engineering. Kitchenham and Charters (2007) provide foundational guidelines for conducting systematic literature reviews, adapted from medical research practices to address the specific challenges of software engineering. These guidelines cover the three phases of a review: planning, conducting, and reporting.

Petersen et al. (2008) introduce systematic mapping studies as a distinct approach from systematic literature reviews. While a systematic review aims to answer specific research questions through in-depth analysis of selected studies, a mapping study aims to produce a broad overview of a research area by classifying publications according to predefined attributes

and identifying patterns, frequencies, and gaps. Petersen et al. (2008) note that systematic mapping is particularly appropriate when a field is large and heterogeneous, precisely the conditions that characterize the sign language dataset landscape. Petersen et al. (2015) subsequently provide revised guidelines that update and clarify the original framework.

2.8.4 Positioning This Thesis

The related work reviewed above identifies a persistent gap: although sign language datasets are widely used and their limitations frequently acknowledged, no existing study provides a structured, systematic mapping of dataset characteristics across the field. Surveys treat datasets as supporting material, dataset papers emphasize individual contributions, and ethics-oriented analyses, such as Bragg et al. (2021), while valuable, do not employ a systematic mapping methodology. Structured review studies, such as Wadhawan and Kumar (2021), classify methods rather than datasets, leaving dataset characteristics without a dedicated structured treatment. This thesis addresses that gap by applying the guidelines of Petersen et al. (2008, 2015) and Kitchenham and Charters (2007) to produce a structured overview of sign language datasets classified by language coverage, modality, scale, annotation, and accessibility.

3. Methodology

This chapter outlines the methodology used to identify, select, analyze, and classify sign language datasets in this study. The approach follows the systematic mapping method proposed by Petersen et al. (2008) and refined in Petersen et al. (2015), supplemented by the systematic review guidelines of Kitchenham and Charters (2007). Each section below corresponds to one of the six steps of the systematic mapping process as defined by Petersen et al. (2008).

3.1 The Systematic Mapping Method

A systematic mapping study is a form of secondary research that produces a structured overview of a research area by classifying existing publications or artifacts according to predefined attributes. The method was introduced by Petersen et al. (2008) and refined in Petersen et al. (2015), drawing on the systematic review tradition established in medical research and adapted for software engineering by Kitchenham and Charters (2007).

The primary aim of a systematic mapping study is to establish the structure of a research field rather than to synthesize quantitative findings, which is the purpose of a systematic literature review. Systematic mapping identifies which topics have been studied, their frequency, the methods used, and existing research gaps. Petersen et al. (2008) describe the output as a visual or tabular map that distributes publications across classification categories derived from the research questions.

3.1.1 Systematic Mapping vs. Systematic Literature Review

Petersen et al. (2008) explicitly distinguish between systematic mapping studies and systematic literature reviews, as these approaches are frequently conflated. A systematic literature review (SLR) is a detailed and focused investigation that aggregates and synthesizes evidence to answer specific research questions, such as "which machine learning algorithm achieves the highest accuracy for isolated sign recognition?" In contrast, a systematic mapping study adopts a broader, classification-oriented perspective, addressing questions such as "what types of datasets exist, and how are they distributed across languages, modalities, and tasks?"

According to Petersen et al. (2015), systematic mapping is most suitable when the research field is large and heterogeneous, when the objective is to identify trends and gaps rather than to synthesize quantitative evidence, and when the scope and structure of the field have not yet

been comprehensively characterized. These conditions are reflected in this study: the sign language dataset landscape encompasses numerous languages, modalities, and task types; the primary aim is classification and gap identification rather than performance synthesis; and the existing surveys reviewed in Section 2.8 address datasets as supporting material rather than as the primary object of structured analysis.

3.1.2 The Six-Step Process

Petersen et al. (2008) define systematic mapping as a six-step process. Each step is described briefly below and is addressed in full in Sections 3.2 through 3.7 of this chapter.

Step 1 — Define research questions. The process begins with the formulation of research questions that determine the scope of the mapping and guide the classification scheme. In contrast to the typically narrow and empirical questions of SLRs, mapping research questions are broad and classificatory, such as "what types of X exist?" or "how is the field distributed across Y?" Petersen et al. (2008) recommend that research questions correspond directly to the classification categories used in the subsequent data extraction step.

Step 2 — Conduct the systematic search. A keyword-based search is performed across multiple academic databases to ensure comprehensive coverage. Petersen et al. (2008) recommend combining several search terms with Boolean operators and searching at least two to three major databases. The search process should be documented in sufficient detail to enable replication. Backward snowballing, which involves examining reference lists of included papers to identify additional relevant sources, is also recommended as a complementary strategy.

Step 3 — Screen publications. All publications retrieved through the search are screened using predefined inclusion and exclusion criteria. Petersen et al. (2008) specify a two-stage screening process: an initial screen based on title and abstract, followed by a full-text review of the remaining records. The same criteria must be applied at both stages, and these criteria should be defined prior to screening to prevent post-hoc adjustments that could introduce bias. Kitchenham and Charters (2007) further emphasize the importance of transparently reporting the number of records retained and excluded at each stage.

Step 4 — Extract data using a predefined keywording scheme. In contrast to systematic reviews, which extract quantitative results, systematic mapping studies focus on extracting classification-relevant metadata from each included publication. Petersen et al. (2008) refer to this as "keywording," which involves identifying the attributes of each study that are relevant to the research questions. A structured extraction form or spreadsheet is used to ensure consistent capture of attributes across all included studies.

Step 5 — Classify and map the results. The extracted data are used to classify publications across the dimensions identified in Step 1. Petersen et al. (2008) present classification results as frequency counts and visual maps that illustrate the distribution of studies across categories, revealing the structure of the research area and facilitating the identification of patterns and concentrations.

Step 6 — Report findings. The final step is to report the classification results, interpret the patterns identified in the map, and explicitly identify research gaps. Petersen et al. (2015) emphasize the importance of discussing both existing coverage and gaps, as gap identification is a primary contribution of a mapping study.

3.1.3 Rationale for Choosing Systematic Mapping

The selection of a systematic mapping methodology, as opposed to a narrative or systematic literature review, was determined by the specific characteristics of the research area and the objectives of the study.

A narrative review is unsuitable in this context because it lacks a structured search process and predefined selection criteria, which increases susceptibility to selection bias and limits replicability. A systematic literature review would require a narrower research question and would aim to synthesize performance results across datasets. However, this approach is currently infeasible due to significant differences in evaluation protocols, metrics, and sign languages across studies, which prevent meaningful quantitative synthesis.

Systematic mapping is the most appropriate approach due to the field's size, heterogeneity, and the lack of prior structured overviews. Petersen et al. (2015) note that mapping is especially suitable when understanding the scope and structure of existing work is necessary before conducting more focused investigations. This context applies to sign language dataset research, where the absence of a structured overview has been identified as an impediment to progress.

3.2 Step 1: Research Questions

Following Step 1 of the Petersen et al. (2008) process, the research questions were formulated to be broad and classificatory, aligned with the mapping objective. The following overarching question and sub-questions were established prior to the commencement of the search:

Research Question: How are publicly available sign language datasets designed, and what research directions do their characteristics enable or constrain?

Sub-question 1: What structural and technical characteristics define existing sign language datasets (such as language coverage, modality, scale, and annotation type)?

Sub-question 2: What types of computational tasks (such as isolated recognition, continuous recognition, translation, and sign spotting) are supported by these datasets?

Sub-question 3: What patterns, limitations, and research gaps emerge from the current dataset landscape?

These questions align with the classification dimensions in the extraction scheme (Section 3.5) and the structure of the results in Chapter 4. Sub-questions 1 and 2 inform the classification categories, while Sub-question 3 informs the gap analysis.

3.3 Step 2: Search Strategy

3.3.1 Data Sources

The search was conducted across five primary academic databases: Google Scholar, IEEE Xplore, ACM Digital Library, SpringerLink, and ScienceDirect. These databases were selected for their comprehensive coverage of computer science, computer vision, and natural language processing research, which are the main disciplines in which sign language datasets are developed and published. In addition, project websites, institutional dataset repositories, and research group pages were searched manually to identify datasets documented outside traditional academic venues.

3.3.2 Search Procedure

A keyword-based search strategy was implemented using an iterative, progressively narrowing approach. The search started with broad umbrella terms to capture a wide perspective and identify the most recognized and frequently cited datasets. The scope was then incrementally narrowed along four dimensions: geographic region, individual country, dataset type, and computational task.

For geographic narrowing, searches moved from continental terms to country-specific terms and then to country-specific acronyms. For example, beginning with "European sign language dataset," the search would narrow to "Finnish sign language dataset" and then to "FinSL dataset" and "FinSL corpus." This approach was applied to all regions and countries represented in the search. For the dataset type, the terms "isolated sign language dataset" and "continuous sign language dataset" were used. For task type, the search used terms such as "sign language recognition dataset," "sign language translation dataset," "sign language

recognition corpus," and "sign language translation corpus." Throughout all iterations, both long-form and abbreviated variants were used, such as "sign language dataset" and "SL dataset," or "sign language corpus" and "SL corpus."

The search terms were combined using Boolean operators to construct queries suited to each database's indexing requirements. The general structure applied was: (*"sign language" OR "SL"*) AND (*"dataset" OR "corpus"*) AND (*"isolated" OR "continuous" OR "recognition" OR "translation"*). This structure accommodates both long-form and abbreviated variants used throughout the search and captures the two primary dataset types and two primary computational tasks targeted by this study.

For each search query, results were reviewed up to the second page. Results beyond page two were examined only when the title or URL contained terms such as "dataset," "corpus," "recognition," or "translation."

The search was supplemented by backward snowballing, which involved reviewing the related work sections and reference tables of all included papers to identify additional relevant datasets. This method was particularly effective for locating datasets not identified by the primary keyword search. Snowballing concluded when five consecutive papers produced no new eligible datasets, consistent with the saturation criterion recommended by Petersen et al. (2008).

3.4 Step 3: Screening

3.4.1 Screening Procedure

The screening process followed the two-stage procedure specified by Petersen et al. (2008) and Kitchenham and Charters (2007). In the first stage, titles and abstracts of all retrieved records were reviewed for the presence of keywords such as "dataset," "corpus," "recognition," and "translation." Records that clearly described gesture vocabularies, synthetic signing data, or were unrelated to any natural sign language were excluded at this stage.

In the second stage, the full texts of the remaining records were evaluated against all inclusion and exclusion criteria specified in Sections 3.4.2 and 3.4.3. Typically, each search iteration produced 10 to 15 candidate records at the title and abstract stage, with 2 to 4 retained after full-text screening. This ratio demonstrates the breadth of the initial search and the specificity of the inclusion criteria. Records that lacked sufficient documentation for attribute extraction or failed to meet the inclusion criteria upon full review were excluded. Figure 1 presents the overall selection process across all three stages.

3.4.2 Inclusion Criteria

A dataset was included if it:

- described a dataset related to a natural sign language, that is, a language used by a deaf or hard-of-hearing community rather than a designed gesture vocabulary;
- provided visual data in the form of video or image recordings;
- was documented in a peer-reviewed paper, conference proceeding, or official dataset description; and
- was documented in English.

3.4.3 Exclusion Criteria

A dataset was excluded if it:

- focused solely on gesture recognition unrelated to any natural sign language;
- lacked sufficient documentation to extract the required attributes;
- consisted solely of synthetic or computer-generated signing data; or
- was not publicly accessible or identifiable through public documentation.

These criteria were defined in the protocol before screening began and applied uniformly across both screening stages, in accordance with the requirements of Petersen et al. (2008) and Kitchenham and Charters (2007).

3.5 Step 4: Data Extraction

3.5.1 Attribute Selection

Consistent with Step 4 of the Petersen et al. (2008) process, a predefined extraction scheme was established before data collection. An initial list of approximately 50 to 60 candidate attributes was compiled based on the research questions and anticipated relevance to researchers. This list was iteratively refined through structured consultation with generative AI tools, resulting in a final set of 38 attributes: 35 core descriptive attributes and 3 supplementary fields for unique features, limitations, and notes.

The 35 core attributes are: Dataset Name, Language(s), Country/Region, Year Published, Authors, Institutions/Affiliations, Publicly Available, License Type, Citation, Dataset Download Link, Number of Signs/Classes, Number of Videos/Samples, Total Hours, Dataset Type, RGB

Video, Depth Video, Skeleton/Pose Data, Multi-view Cameras, Audio Included, Recording Environment, Background Complexity, Lighting Conditions, Camera Type, Resolution, Frame Rate, Signer-Independent Split, Baseline Model Results, Gloss Labels, Sentence Alignment, Non-Manual Annotation, Number of Signers, Deaf/Hearing/CODA Representation, Gender Distribution, Age Range, and Dialect/Regional Variation.

3.5.2 Extraction Process

Data extraction proceeded in two stages. First, attribute values were recorded exactly as stated in the source documentation, without interpretation. Second, the extracted data were standardized to ensure uniform formatting across all 30 datasets. Standardization included removing extraneous white space, normalizing variant expressions of the same value into a single canonical form (such as consolidating "Not explicitly stated," "Not mentioned," and "Not explicitly listed" to "Not specified"), correcting formatting inconsistencies, and ensuring consistent use of units and parentheses. When documentation was ambiguous or missing, the value was recorded as "Not specified" to maintain accuracy and avoid bias, in line with Kitchenham and Charters (2007).

3.6 Step 5: Classification and Mapping

3.6.1 Classification Dimensions

Following Step 5 of the Petersen et al. (2008) process, the extracted attributes were organized into seven thematic groups: General Information, Scale and Scope, Modalities, Recording Conditions, Annotation and Benchmarking, Signer Demographics, and Remarks. These groups served as the basis for the classification and mapping of all 30 datasets.

The classification addresses the following six dimensions, which correspond directly to the research questions formulated in Section 3.2:

- language coverage and geographic distribution (Sub-questions 1 and 3);
- dataset type and supported computational tasks (Sub-questions 1 and 2);
- modality and recording setup (Sub-question 1);
- annotation depth and structure (Sub-questions 1 and 2);
- benchmark rigour, including signer-independent splits and baseline results (Sub-questions 1 and 3); and
- demographic reporting and signer diversity (Sub-questions 1 and 3).

3.6.2 Summary Tables

Classification results are presented in Chapter 4 using five structured summary tables, consistent with the tabular map output described by Petersen et al. (2008).

Table 1: Scale and Scope presents dataset name, year, type, number of sign classes, number of video samples, total duration in hours, and number of signers for all 30 datasets, sorted in the original spreadsheet order.

Table 2: Modality Coverage displays the availability of RGB video, depth video, skeleton/pose data, multi-view cameras, and audio for each dataset, with cells color-coded green for Yes, red for No, and amber for Partial.

Table 3: Annotation and Benchmarking Characteristics covers gloss labels, sentence alignment, non-manual annotation, signer-independent split, and baseline results, using the same color-coding scheme as Table 2.

Table 4: Recording Conditions presents recording environment, background complexity, lighting conditions, camera type, resolution, and frame rate for each dataset, with row colors distinguishing studio, natural or online, and mixed or lab settings.

Table 5: Signer Demographics covers the number of signers, deaf/hearing/CODA status, gender distribution, age range, and dialect or regional variation. A completeness summary at the bottom of the table indicates how many datasets report each demographic field.

Each table is followed by a "Note" paragraph that defines column values and clarifies how attributes were extracted, in accordance with standard academic table formatting.

3.6.3 Analytical Spreadsheet

In addition to the summary tables, a structured analytical spreadsheet was developed to support the classification and gap analysis. The spreadsheet contains eight analytical sheets beyond the main data and cleaned-up data sheets:

- **Dashboard:** provides aggregate statistics including total datasets, unique sign languages, combined duration, median signers, median sign classes, and annotation coverage rates, as well as type distribution, regional breakdown, modality coverage, annotation feature coverage, the ten largest datasets by duration, and the full list of sign languages covered.
- **Overview:** provides one row per dataset with the 14 most frequently referenced columns: name, language, country, year, type, duration, sign classes, video samples, signers, public availability, license, download link, and citation.

- **Scale and Scope:** sorts all 30 datasets by duration and presents both raw text values and extracted numeric values for signs, videos, hours, and signers, with a summary statistics block reporting the median, minimum, and maximum for each numeric dimension.
- **Modalities:** displays RGB, depth, skeleton/pose, multi-view, and audio availability per dataset with color-coded cells, and includes a summary count of Yes, No, and Partial responses across all 30 datasets for each modality.
- **Annotation and Benchmarking:** sorts datasets by their annotation score (0 to 5), which accumulates one point each for gloss labels, sentence alignment, non-manual annotation, a signer-independent split, and baseline results, with score cells color-coded green for 4 or above, amber for 2.5 or above, and red for below 2.5.
- **Geographic and Linguistic Coverage:** groups datasets by region with region-coded row colors, with a summary listing the five regions and the datasets belonging to each.
- **Signer Demographics:** sorts datasets by number of signers and presents raw and numeric signer counts, deaf/hearing/CODA representation, gender distribution, age range, and dialect variation, with a completeness table showing how many datasets report each demographic field.
- **Recording Conditions:** covers recording environment, background complexity, lighting, camera type, resolution, and frame rate, with row colors distinguishing studio, natural or online, and mixed or lab settings.

3.7 Validity Considerations

Petersen et al. (2008) identify several threats to validity that are common in systematic mapping studies, and Kitchenham and Charters (2007) provide corresponding mitigation strategies. The following threats and mitigations apply to this study.

Search for incompleteness. Relevant datasets may be absent from the selected databases, documented only in non-English sources, or distributed through informal channels. This risk was mitigated by searching five academic databases in parallel, supplementing the keyword search with manual searches of project websites and institutional repositories, and applying backward snowballing until saturation was reached.

Selection bias. Subjective judgements during title and abstract screening may introduce inconsistency. The inclusion and exclusion criteria were defined in the protocol before screening began and applied uniformly across both screening stages.

Extraction inconsistency. Ambiguous or incomplete dataset documentation may lead to inconsistent attribute recording. All ambiguous values were recorded as "Not specified" rather than estimated, and the same extraction scheme was applied to every dataset.

Classification subjectivity. The assignment of datasets to classification categories involves judgement. Classification dimensions were defined before extraction began and are directly derived from the research questions, reducing the scope for post-hoc adjustment.

These mitigation strategies are consistent with the recommendations of Petersen et al. (2008, 2015) and Kitchenham and Charters (2007), and are documented to enhance the reproducibility and transparency of this study.

4. Results and Initial Discussion

This chapter presents the results of the systematic mapping study described in Chapter 3. Applying the Petersen et al. (2008) framework, nine sign language datasets were identified, selected, and analysed according to the predefined extraction scheme. The findings are organised around the six classification dimensions introduced in Section 3.7: language coverage, dataset type, scale, modality, annotation depth, and benchmark and accessibility practices.

4.1 Overview of Included Datasets

A total of seven sign language datasets were identified through the systematic search and selection process outlined in Chapter 3. These datasets were chosen to ensure coverage across various sign languages, dataset sizes, and design features. The selected datasets cover sign languages from several geographic regions and differ in dataset type, modality, annotation detail, and accessibility. Each dataset is described in peer-reviewed publications and includes sufficient metadata for structured analysis. This chapter reports the results of dataset characterization according to the predefined extraction criteria. The findings are structured around key dimensions relevant to dataset usability, accessibility, and research applicability.

4.2 Distribution by Sign Language and Region

The seven datasets cover a diverse range of sign languages, including national sign languages from Europe, South America, Asia, and North America. Each dataset targets a single sign language, reflecting the language-specific approach to sign language data collection. Despite broad geographic representation, the distribution is uneven. Most originate from regions with

established research infrastructure in computer vision and machine learning. Several regions and sign languages are underrepresented, indicating gaps in dataset availability. These findings highlight the fragmented development of sign language datasets, which are typically created within isolated national or institutional contexts rather than through coordinated international initiatives.

4.3 Dataset Type and Task Support

Analysis of dataset types indicates that most included datasets focus on isolated sign recognition. These datasets generally consist of short video clips, each representing a single sign. Few datasets support continuous signing, which involves sequences of signs forming sentences or phrases. Continuous datasets are more complex to collect and annotate, which may explain their lower prevalence. As a result, current publicly available datasets are more suitable for isolated sign recognition tasks than for advanced applications such as sign language translation or discourse-level analysis.

4.4 Dataset Scale and Composition

The included datasets vary substantially in scale. Some contain a limited number of sign classes and samples, while others include hundreds of signs and tens of thousands of video instances. The number of signers also varies considerably across datasets. Larger datasets typically include more signers, supporting greater variation in signing style and appearance. However, detailed demographic information, such as age range or linguistic background, is frequently absent or inconsistently reported. Variation in scale and documentation affects direct comparison between datasets and restricts the assessment of signer diversity across studies.

4.5 Modalities and Recording Setup

Most datasets utilize RGB video as the primary modality. Several datasets also provide additional modalities, such as depth data or skeletal pose information, typically captured with depth sensors. Multimodal datasets provide richer representations but are less common and generally require specialized hardware. In contrast, RGB-only datasets are more prevalent and easier to collect, though they may offer limited detail for fine-grained sign analysis. Recording environments range from controlled studio settings with simple backgrounds to more complex environments with varied backgrounds and lighting. Controlled setups are predominant in

isolated-sign datasets, whereas continuous datasets typically exhibit greater environmental variability.

4.6 Annotation and Benchmarking Practices

All included datasets offer some form of annotation, most commonly gloss-level labels for individual signs. However, annotation depth and structure differ considerably across datasets. Only a subset of datasets explicitly defines signer-independent evaluation protocols, in which test data consists of signers not present during training. These protocols are essential for evaluating model generalizability but are not consistently implemented. Not all datasets provide baseline model results. When included, baseline evaluations vary in model architecture, evaluation metrics, and reporting standards, which limits direct performance comparisons across datasets.

4.7 Accessibility and Availability

All datasets included in this initial analysis are described as publicly available or accessible upon request for research purposes. However, licensing information is often incomplete or vaguely specified. Documentation quality also varies. Some datasets include comprehensive descriptions, download instructions, and evaluation protocols, whereas others provide minimal guidance beyond the associated publication. These differences affect dataset accessibility and practical usability, especially for researchers who are new to the field.

4.8 Initial Discussion

The initial mapping indicates a strong emphasis on isolated sign recognition datasets and controlled recording environments. While this supports early-stage model development, it restricts the evaluation of model generalization and real-world applicability. The lack of consistent signer-independent benchmarks and incomplete demographic reporting further complicates cross-dataset comparisons. These findings indicate a need for improved dataset documentation practices and standardized evaluation criteria, which will be explored further as the dataset corpus expands.

5. Conclusion

This thesis presents a systematic mapping and initial benchmarking of sign language datasets, offering a structured overview of their key characteristics and current limitations. By identifying and categorizing datasets based on language coverage, modality, scale, and annotation, this work establishes a clear baseline for understanding the current landscape of sign language research datasets. The initial findings indicate that many available datasets primarily support isolated sign recognition and are collected in controlled environments. Although these datasets are valuable for early-stage research, they provide limited support for evaluating generalization and real-world applicability. Furthermore, inconsistencies in dataset documentation and evaluation practices hinder direct comparison across datasets. The mapping and benchmarking framework proposed in this thesis serves as a foundation that can be extended as additional datasets are incorporated. Future research can build on this baseline by expanding the dataset corpus, refining evaluation criteria, and examining inclusivity and accessibility in greater detail. Overall, this study provides a transparent starting point for more systematic analysis and informed utilization of sign language datasets.

References

1. Bragg, D., Caselli, N., Hochgesang, J. A., Huenerfauth, M., Katz-Hernandez, L., Koller, O., Kushalnagar, R., Vogler, C., and Ladner, R. E. (2021). The FATE landscape of sign language AI datasets: An interdisciplinary perspective. *ACM Transactions on Accessible Computing*, 14(2), Article 7. <https://doi.org/10.1145/3436996>
2. Koller, O. (2020). *Quantitative survey of the state of the art in sign language recognition* (arXiv:2008.09918). arXiv. <https://arxiv.org/abs/2008.09918>
3. Camgöz, N. C., Hadfield, S., Koller, O., Ney, H., and Bowden, R. (2018). Neural sign language translation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 7784–7793. <https://doi.org/10.1109/CVPR.2018.00812>
4. Camgöz, N. C., Koller, O., Hadfield, S., and Bowden, R. (2020). Sign language transformers: Joint end-to-end sign language recognition and translation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 10023–10033. <https://doi.org/10.1109/CVPR42600.2020.01004>
5. Forster, J., Schmidt, C., Koller, O., Bellgardt, M., and Ney, H. (2014). Extensions of the sign language recognition and translation corpus RWTH-PHOENIX-Weather. In *Proceedings of the 9th International Conference on Language Resources and Evaluation (LREC 2014)*, pp. 1911–1916. <https://aclanthology.org/L14-1472/>
6. Sarhan, N. and Frintrop, S. (2024). A review of deep learning-based approaches to sign language processing. *Advanced Robotics*, 39(1–2), 1–20. <https://doi.org/10.1080/01691864.2024.2442721>
7. Duarte, A., Palaskar, S., Ventura, L., Ghadiyaram, D., DeHaan, K., Metze, F., Torres, J., and Giro-i-Nieto, X. (2021). How2Sign: A large-scale multimodal dataset for continuous American Sign Language. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2735–2744. <https://doi.org/10.48550/arXiv.2008.08143>
8. Li, D., Rodriguez, C., Yu, X., and Li, H. (2020). Word-level deep sign language recognition from video: A new large-scale dataset and methods comparison. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pp. 1459–1469. <https://doi.org/10.48550/arXiv.1910.11006>
9. Petersen, K., Feldt, R., Mujtaba, S., and Mattsson, M. (2008). Systematic mapping studies in software engineering. In *Proceedings of the 12th International Conference on Evaluation and Assessment in Software Engineering (EASE '08)*, pp. 68–77. <https://doi.org/10.14236/ewic/EASE2008.8>
10. Petersen, K., Vakkalanka, S., and Kuzniarz, L. (2015). Guidelines for conducting

- systematic mapping studies in software engineering: An update. *Information and Software Technology*, 64, 1–18. <https://doi.org/10.1016/j.infsof.2015.03.007>
11. Rastgoo, R., Kiani, K., and Escalera, S. (2021). Sign language recognition: A deep survey. *Expert Systems with Applications*, 164, 113794. <https://doi.org/10.1016/j.eswa.2020.113794>
 12. Sincan, O. M., & Keles, H. Y. (2020). AUTSL: A large scale multi-modal Turkish Sign Language dataset and baseline methods. *IEEE Access*, 8, 181340–181355. <https://doi.org/10.48550/arXiv.2008.00932>
 13. Wadhawan, A. and Kumar, P. (2021). Sign language recognition systems: A decade systematic literature review. *Archives of Computational Methods in Engineering*, 28, 785–813. <https://doi.org/10.1007/s11831-019-09384-2>
 14. Sarhan, N. and Frintrop, S. (2023). Unraveling a decade: A comprehensive survey on isolated sign language recognition. In *Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops (ICCVW)*, pp. 3210–3219. <https://ieeexplore.ieee.org/document/10350553>
 15. Tan, S., Khan, N., An, Z., Ando, Y., Kawakami, R., and Nakadai, K. (2024). A review of deep learning-based approaches to sign language processing. *Advanced Robotics*, 38(23), 1649–1667. <https://doi.org/10.1080/01691864.2024.2442721>